

Advice, hints, and general information from the jury

- Your submissions will be run multiple times, on several different input files. The order in which the input files are being used is fixed and does not change from one submission to another.

If your submission is incorrect, the error message you get will be the error that occurred on the first input file on which you failed. E.g., if your instance is prone to crash but also incorrect, your submission may be judged as either “Wrong Answer” or “Run Time Error”, depending on which is discovered first.

- Your submission will have exactly one processor core fully at its disposal while running. You should not attempt to use multithreading in your submissions. Attempts to start extra processes or threads may lead to a judgement of “Run Time Error”.
- Each problem will have a stated time limit. This time limit is per input file and refers to CPU time, which includes time spent on slow parsing (e.g., using `Scanner` in Java).

All judging machines are physically identical, running on an Intel Core i5-4570 CPU @ 3.20GHz. See <https://www.nwerc.eu/system/> for the exact compiler versions and command line arguments used to run your submissions, and <https://www.nwerc.eu/vm/> for a virtual machine that includes the exact same compiler versions as the judging machines.

- The memory limit is the same for all problems and is 2048 MB.
- For problems with large I/O, it is in your best interest to use buffered input and output in Java. We have tried to make sure that the I/O-intensive problems are solvable using even slow I/O. But for some problems it can be more difficult to get accepted using e.g. the `Scanner` class, since it consumes a significant amount of valuable CPU cycles that could otherwise be used by your algorithm. Also consider explicit output buffering (e.g. by using the `BufferedWriter` class) to avoid flushing every line of output separately.
- We guarantee that every problem is solvable in C, C++, Java and Python 3. There is no such guarantee for Kotlin.
- This year the contest includes interactive problems. Each interactive problem is accompanied with a Python script to help with testing your solutions locally. The tool attempts to detect and report common errors, but it is not guaranteed that a program that passes the testing tool will be accepted.

The tool can be found in the Kattis interface, in the panel below the statement of the corresponding problem, under **Downloads > testing_tool.py**. The source code of the tool contains instructions on its use. If you haven’t already, make sure you install **Python** before the start of the contest. You will be able to test this during the practice session.

- In solutions to interactive problems, make sure to flush after every write to standard output. Failing to do so may result in the interactor waiting for input indefinitely, and your solution being judged as “Time Limit Exceeded”.
- Verdicts for incorrect submissions to interactive problems may not be deterministic. We can guarantee the following; a verdict of “Wrong Answer” means that your submission printed something wrong, and a verdict of “Run Time Error” means that your submission returned a non-zero exit code. If your submission does both, you may get either verdict.
- Lowercase letters are the 26 letters “a” to “z”. Uppercase letters are the 26 letters “A” to “Z”.
- We generally try to make the input and output formats as easy as possible to follow for you as contestants — input and output are meant to be the easiest parts of the problems!

- For input, multiple numbers or words on one input line will be separated by single spaces. There are no additional spaces, tabs or newlines.
- For output, unless explicitly stated in the problem statement we do not require you to output exactly the right number of spaces and newlines. Errors in whitespace within reason will be accepted, but if all whitespace between two tokens are omitted, that is (of course) an error.
- For problems with real-valued output, we generally only require that your output is correct up to a certain either relative or absolute error. For example, if the problem statement requests a relative or absolute error of at most 10^{-6} , this means that:

- * If the correct answer is 0.005, any answer between 0.004999 and 0.005001 will be accepted.
- * If the correct answer is 500, any answer between 499.9995 and 500.0005 will be accepted.

In these cases, any reasonable format for floating point numbers is acceptable. For instance, “17.000000”, “0.17e2”, and “17” are all acceptable ways of formatting the number 17. For the definition of reasonable, please use your common sense.